

Radio Units

Cell 0

Cell 1

Cell 2

Cell 3

*An Uplink Slot
Every 2.5 ms*

Samples

One GPU of the AI-RAN Server

Conventional Receiver

Decodes Every TB. Result Due in 4.0 ms

Neural Receiver (NeuralRx)

Recovers Failed TBs. Result Due in 11.5 ms

AI Inference

Serves LLM Requests. Answer Due in 200 ms

Data

*Retransmits
What Failed*

Layer 2



AI Requests

53k tokens/s Offered