


GPU Share for AI

Low Priority, 30% Cap

NeuralRx

AI

AI Is Always on the GPU, Also Next to the NeuralRx



AI Served: **18.0k** tokens/s

Recoveries Lost: **1.2%**

Antiphase

This Work

NeuralRx

AI

AI Stops While the NeuralRx Runs, and Has the Whole GPU Otherwise



AI Served: **34.1k** tokens/s

Recoveries Lost: **0.1%**

1.9× the AI, 12× Fewer Lost