


## AI Requests



*Each Due  
in 200 ms*



### Request Dispatcher

Predicts the Finish Time  
on Every GPU From Its  
Queue and Recent Rate

*Earliest  
Finish*



**GPU 0**




**GPU 1**



*Shortest Queue*

**GPU 2**



**GPU 3**



*Queued AI Work on Each GPU*



Not Admitted: Cannot Finish in 150 ms

*75% of the 200 ms Limit. The Same Dispatch  
and Admission for Every Compared Method*