



Server: Four A100 GPUs, 16 Cells of 100 MHz



*AI: LLM Requests (Qwen2.5-1.5B), up to 1,024 Tokens,
53k tokens/s Offered, Each Due in 200 ms*

Three Radio Load Patterns

(a) Steady Full Load



(b) Load Changes Every 2 s



(c) Random Load Steps



Active Cells Over Time