



one GPU: on-time
AI tokens per second

